

3D Gaussian Splatting for real time

Radiance fields Rendering

- traditional 3D scene (meshes, point clouds) .
- NeRF → extensive training & rendering

3D Gaussian Represent

- ↳ preserves volumetric radiance fields prop.
- ↳ Enable efficient training.
- ↳ Real time - rendering at 30 FPS.

Key Innovation

- ↳ 3D gaussian repres. derived from SGM points.
- ↳ Adaptive density control & anisotropic Gaussian optimization
- ↳ Fast tile based splatting renderer for R.T. Rendering.

Point Based Rendering (PBR) → 3D objects are rendered as a points.

directly rendering point samples without requiring explicit inform. like edges or faces. rather than triangles or volumetric grid (voxel)

(depth maps, Lidar).

• Can result in holes or aliasing

• Harder to process

• need specialized technique like Splatting for smooth rendering

point cloud

→ surface

3D pos (x, y, z)

- color (R, G, B)

- Normal vector (optional)

- opacity (α)

- Spherical Harmonics (SH) Coefficient (per view dependent effects)

These points can be obtained from

- GFM, MVS, LiDAR, NeRF,

point rendering technique.

→ point sprites → each 3D point is treated as 2D texture (simple & fast) but suffer from holes at sparse region (sprite)

→ splat based rendering

uses elliptical splats (disc) instead of single points.

Blends overlapping splats to fill gaps and create a smooth appearance.

→ deferred shading.

→ First renders point attributes into a buffer.

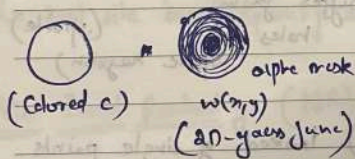


→ Then applies shading & lighting in a second phase.

↳ Adaptive Density Control

Visibility and occlusion handling

↳ Since points don't have explicit connectivity, rendering requires depth based techniques.

- ① Depth Sorting → sort points based on their depth for correct rendering
- ② α -Blending (opacity Blending) → used in GIS for smooth rendering
- ③ Screen space culling → removes points outside the view frustum.



Point Based Rendering Eqⁿ

$$C = \sum_{i=1}^N T_i d_i C_i \quad \alpha_i = 1 - \exp(-\omega_i d_i) \text{ is the opacity of each point.}$$

Blending accumulates color along the ray,

Similar to Nelder's ray marching.

anisotropic Gaussian \rightarrow allowing elliptical splds.
for better approx. Complex Scene

Covariance (Σ) \rightarrow Gated Gaussian shape. Structure

Spherical Harmonics (SH) \rightarrow model color & lighting.

Objective \rightarrow High Quality Novel view synthesis
goals \rightarrow generate new views of a 3D scene from
a sparse set of images.

Input: SfM with no Surface Normals

Challenge: Need a representation that is both
differentiable and fast for rendering.

Why 3D?
Gs. they can be optimized using
gradient descent (differentiable)

- \rightarrow They can be easily projected onto a 2D-plane
- \rightarrow They do not require surface normals, which are hard to estimate from SfM data.

adam optimize for gradient updates.

$$G(\mu) = e^{-\frac{1}{2}(\mu - \mu)^T \Sigma^{-1}(\mu - \mu)}$$

mean 3D pos of Gaussian

$\Sigma \rightarrow$ Covariance matrix

(define shape & spread)

Why no Surface Normals?

- \rightarrow prev. point based methods used small 2D circles with Normals
- \rightarrow SFM points are too sparse, making it difficult to estimate accurate Normals
- \rightarrow Normals are too Noisy, making optimization unstable.
- \rightarrow 3D G.S. requires geometry without needing normals.

Projection of 3D G. to 2D for Rendering.

Viewing transform matrix V (World Coordinate to Camera Coordinate)

② Covariance Matrix Σ' in Camera Space

$$\Sigma' = J W \Sigma W^T J^T$$

Jacobian Matrix (approx the perspective transform)

by removing third column & row
2x2 (Covar. Matrix)

Optimization the Covar. Matrix Σ

\hookrightarrow without opti, the Gaussian would be isotropic (spherical)
 \hookrightarrow would not adapt to complex geometry of objects.

Covariance Matrix must be +ve semi-definite

↳ direct optim. of C.M is risky due to 6 indep. param.

The paper factorize CM into

① S (Scaling Matrix)

(avoids Numerical instability)

② R (Rotation Matrix)

$E = SAS^T R$ → ensure stability during optimization.

→ Gradient descent (or sometimes produce invalid E matrices, causes rendering artifacts or crashes.

Tile Based Rasterization →

• ~~Sorting~~

• prev. methods had hard limit on the no of splits that could receive gradient descent.

*⁴ Sorting & Blending is Computat. extensive

Tile based rendering → 16×16 tiles ← Screen

(Gaussian Scaled) process per tile

→ cull Gaussian outside the view frustum. (99% Gof)

→ Radix Sort → sort G in depth order. (front to back)
CPU (fast sorting)

Gaussian are blended using (α)

If incorrect sorting occur further C_n may occur in front of closer ones.

Efficient Gradient Computation (Backward pass)

→ prev method store long lists of blended points per pixel,

Instead of storing, the method reuses the sorted gaussian from the forward pass. (Memory efficient)

→ accumulates color until α reaches 1, stop early reducing computation. (Efficient α blending)

Limitations: artifacts in poorly observed region.

Regularization could improve handling of unseen areas.

Training large scene requi > 20 GB GPU M.

↳ point cloud compression technique could reduce memory usage.

Training start from 100K Random Gaussian → 6K-10K meaningful Gauss

↑ PSNR → image quality
↑ SSIM → measure similarity
↓ L2-P2PS → Capture fine details.

Paperkraft

Spherical Harmonics in 3DGS.

→ Instead of storing RGB Glor(c) per Gaussian, SH encodes how the Glor of the Gaussian changes on the viewing direction.

$$C(r) = \sum_{l=0}^L \sum_{m=-l}^l c_{lm} Y_l^m(r)$$

learned SH coeffs of each gaussian
↳ S-H basis func.

↳ Glor as a func of viewing dir. (r)

provides efficient way to store directional appearance

Low SH are fast but smooth out fine details

$$L=0, 2, 2$$

High SH ($L=3+$) allows better reflection but increase memory usage.

GS-IR (Inverse Rendering)

Inverse Rendering \rightarrow reconstruct 3D scene prop from 2D images by estimating geometry, material prop & illumination.

(hemisphere of all possible light dir)

$$L_o(x, u) = \int_{\Omega} \underbrace{L_i(x, l)}_{\text{incoming BRDF}} \underbrace{f_r(l, u)}_{\text{Surface refl}} (l \cdot n) dl \rightarrow \text{rendering eqn.}$$

(out rad at x in direction u)
(incoming BRDF at x in direction l from Ω)
Surface refl

① In forward rendering, we compute $L_o(x, u)$ given scene prop.

② In Inverse Rendering, we start with $L_o(x, u)$ (observed image)

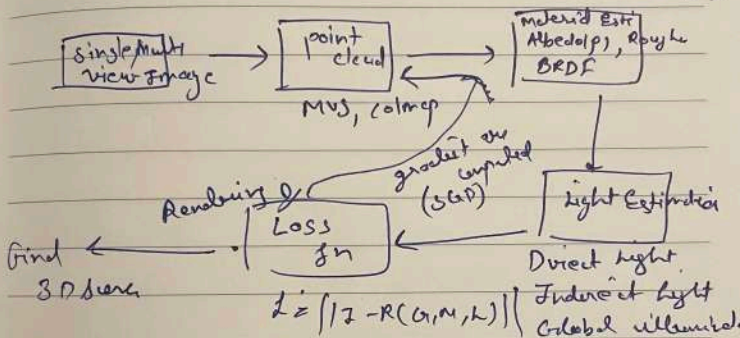
and estimate

- ① geometry
- ② Materials
- ③ lighting.

Inverse Rend is ill posed, multiple combination of G, M, L can produce the same image.

We use differentiable rendering, optimization & DL Tech.

$\hat{G}, \hat{M}, \hat{L} = \arg \min_{G, M, L} \mathcal{L}(\hat{Z}, R(G, M, L))$
 \hat{L} (estimated parameter) loss funcⁿ (diff b/w, R ed & Re-rendered Image)



NeRF \rightarrow Representing ^{Scene} Neural Radiance field for view synthesis

View Synthesis \rightarrow 2D images set \rightarrow new arbitrary angle view

Challenges

- ① Complex Geometry
- ② view dependent lighting
- ③ high resolution details.

Models a 5D func which

- ① Encodes a 3D point (x, y, z) & viewing direction (θ, ϕ) .
- ② outputs :
 - ① Color (r, g, b) Defines what color is seen at that point.
 - ② Density (σ) . How opaque the point is

This func is learned using a MLP (fully Connected) without convolutions layer.

★★

Prior methods use discrete voxels grids or meshes

→ lot of memory to store H.R 3D scenes.

→ Meshes require pre-defined geometry
(which isn't avail for real-world images)

Neural Continuous 3D Representation

→ no explicit geometry is needed

→ neural networks learns the underlying structure of the scene.

Storage efficient → Scene is encoded in the network's weights.

Neural works?


→ render an image from novel viewpoint

① Sample along camera rays
points

② Predict color & density for each point.

Input (x, y, z, θ, ϕ) in MLP

output (r, g, b, σ) values



Scene Representation 5D Continuous 3D spatial location (x, y, z)
2D View Dir (θ, ϕ)

3D coordinates \rightarrow 3 FCL \rightarrow (Volume Density)
(256) Relu (depends on x)
+ 256 Dimensional Feature Vector

Feature Vector + viewing direction \rightarrow FCL \rightarrow RGB
(128c + Relu) (C)
(6br)

Handling view-Dependent Effects

\rightarrow Non-Lambertian Surface

Nerfactors

Shape & reflectance of an object from multi-view
Captured under one unknown illumination

Goal \rightarrow to enable scene free viewpoint relighting,
& material editing, by factorizing scene into
albedo, light visibility & BRDF.

Use Nerf to estimate volumetric geometry, then
distilled into surface representation,

- ① Geometry
- ② Spatially varying Reflectance
- ③ Light visibility
- ④ Environment lighting

Previous work:

NERF \rightarrow require multiple lighting conditions.

NERF doesn't model shadow & light visibility

PhysCN \rightarrow assumes non-spatially-varying reflectance

NeXfactor

- uses learned BRDF instead of analytic BRDF
- works with single unknown illumination.

Method

Input assumption

- Starts with multi-view images of an object & their camera pose, captured under single unknown lighting conditions.

- Representing using Neural Fields.

Each of these fields is parameterized by MLP. The MLP's are trained so that when the model "renders" the object from these neural fields, the output matches the observed input images.

Outputs

for every 2D point x on the object's surface

NeXfactor produces

- ① Surface Normal (n)
- ② Light Visibility ($v(w_i)$)
- ③ Albedo (a)
- ④ Reflectance (2BRDF)